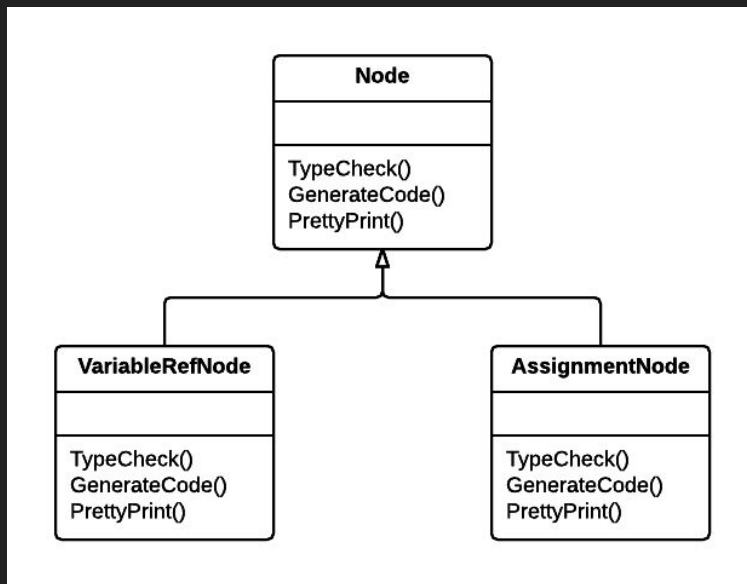


Visitor Pattern

Behavioral, Object focused

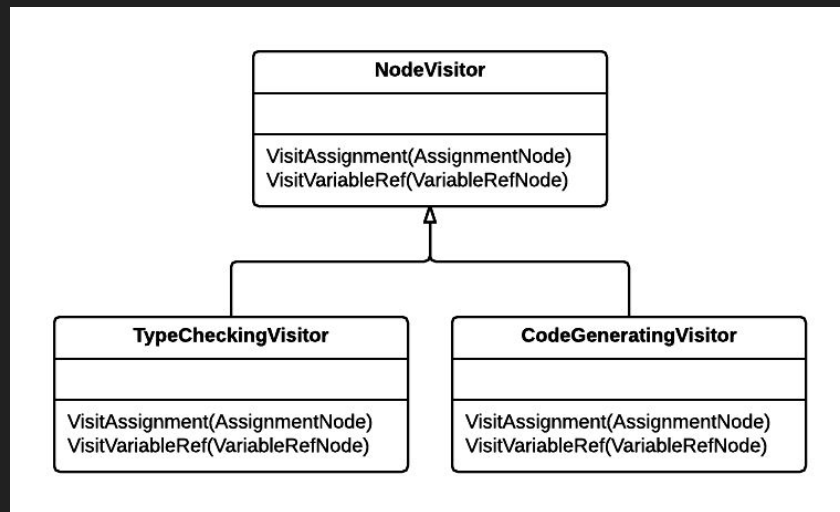
Motivation

- Consider a compiler representing a program as an abstract syntax tree (AST)
- The compiler must perform operations on the AST for “semantic” analysis
 - Type-checking
 - Code optimization
 - Flow analysis
 - Etc.
- This requires a lot of nodes with a lot of operations
- Maintaining this code base quickly becomes a nightmare

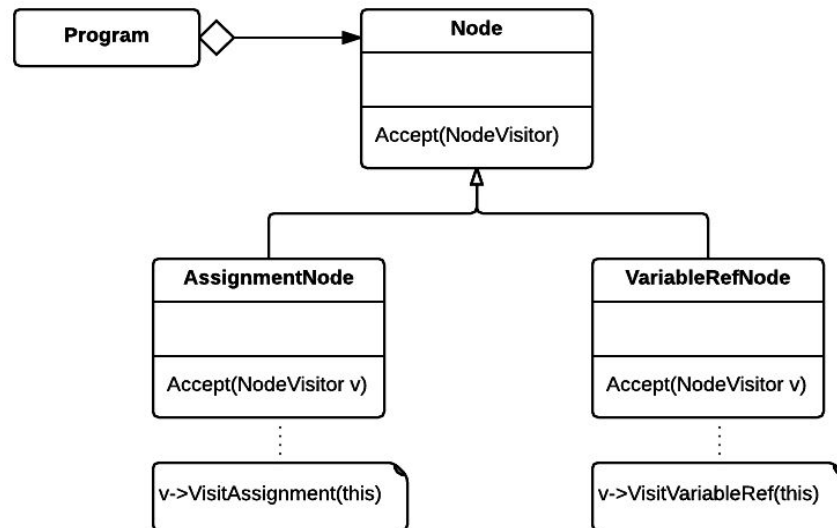
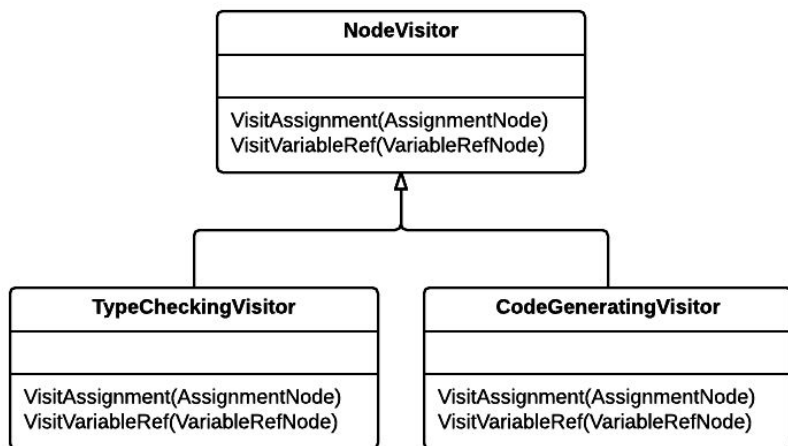


Motivation

- It would be better if new operations could be added separately and node classes were independent of the operations applied to them
- This is obtained by packaging related operations in a separate class called **Visitor**
- These **Visitor** objects are passed to elements of the AST as it is traversed

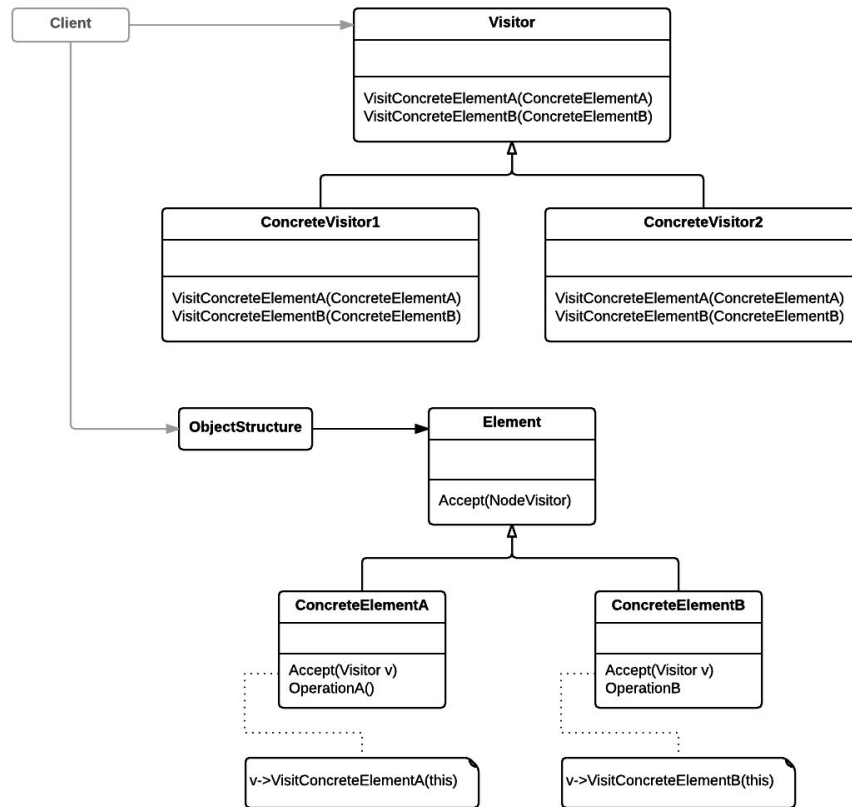


Motivation

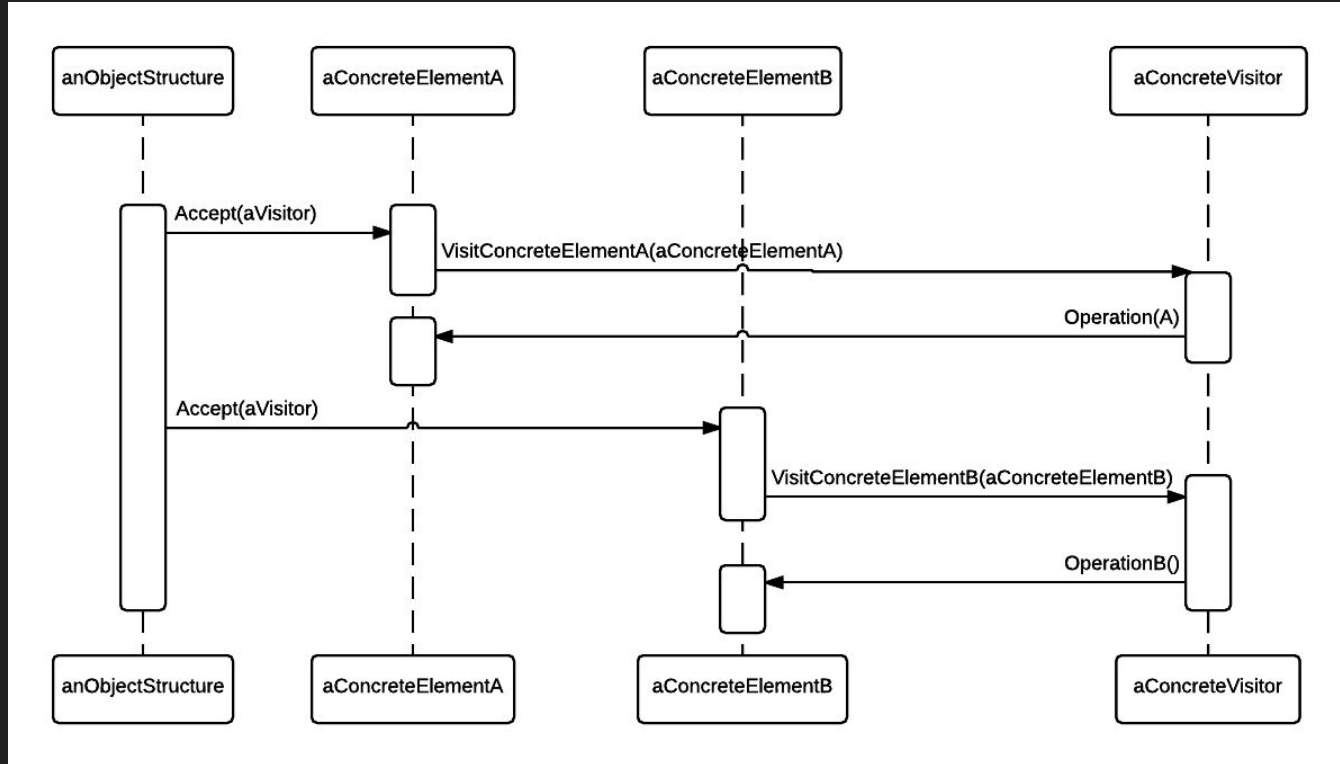


Structure

- **Visitor** - Declares a Visit operation for each class of ConcreteElement in the object structure
- **ConcreteVisitor** - Implements each operation declared by Visitor
- **Element** - Defines an `Accept()` operation with visitor as an argument
- **ConcreteElement** - Implements the `Accept()` operation
- **ObjectStructure** -
 - Can enumerate its elements
 - Can be a composite or collection
 - May provide high level interface



Structure



Consequences

- **Pros:**

- Adding new operations is easy
- Related operations are gathered together, unrelated operations are separated
- Visitors can visit across different class hierarchies

```
Class Visitor {  
    Public:  
        Void VisitMyType (MyType*);  
        Void VisitYourType (YourType*);  
        ...  
}
```

- Visitors can accumulate state while visiting each element

- **Cons:**

- Adding new `ConcreteElement` classes is difficult
- Visitors assume the `ConcreteElement` interface lets them do their job, which typically breaks encapsulation

Example from “Design Patterns: Elements of Reusable Object-Oriented Software”

Erich Gamma

Richard Helm

Ralph Johnson

John Vlissides

Traversal vs. Traversal Actions

- The iterator provided a method of traversing the glyph structure **BUT**, it did not provide a way to check spelling and do the hyphenation

Where do we put the responsibility for analysis?

Traversal vs. Traversal Actions

- Putting the responsibility for analysis in the Iterator makes analysis an integral part of traversal, which it isn't necessarily a part of
- We want to **distinguish** between the traversal and the analysis being performed
 - This provides more flexibility and potential for re-use

So, Where do we put the responsibility for analysis?

Analysis

- A given analysis must be able to distinguish between different kinds of glyphs
 - Spelling and hyphenation only care about visible glyphs
- We don't want to put the analytical capabilities into the glyph classes themselves
- Each analysis could add one or more abstract operations, and each subclass would have to implement them with the role they play

Analysis

- We then have to change every glyph class whenever we add a new analysis, even if it isn't affected!
- This can be partially fixed by providing a default implementation in the base class, however we still need to expand the Glyph interface with each new analytical capability

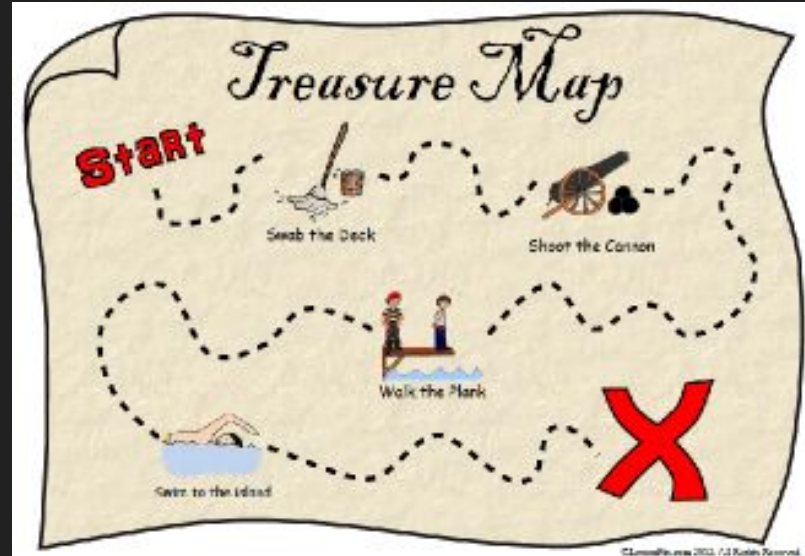


Encapsulating the Analysis

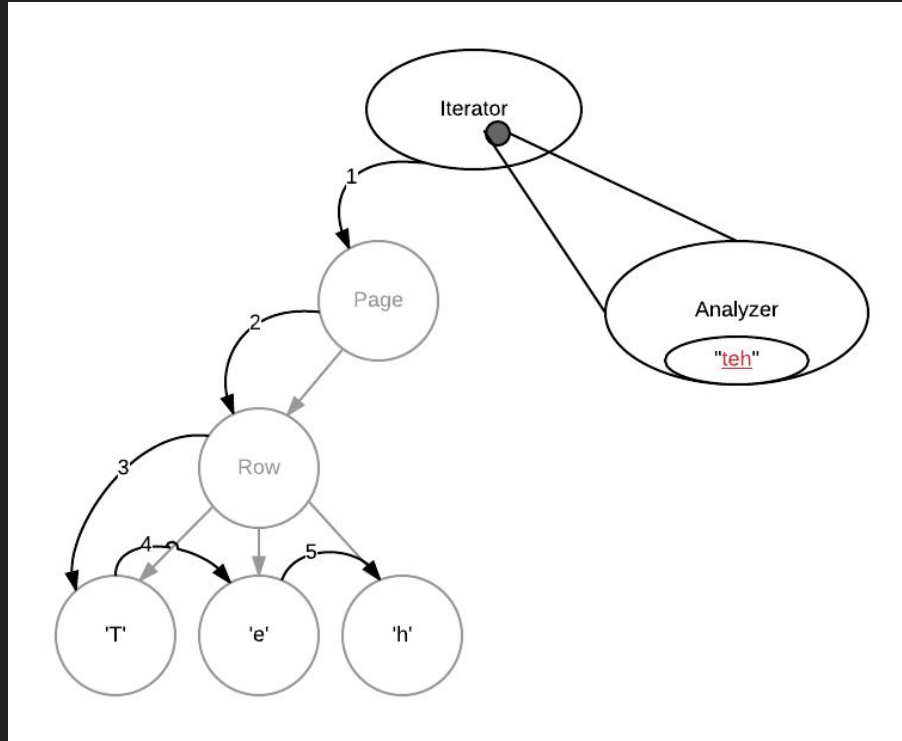
- The analysis should be encapsulated in a separate object
- This analysis object can interact with the appropriate *iterator* to perform its analysis
- The *iterator* object “carries” the glyph instance associated with it
- The **analysis** object is able to visit this instance and perform a piece of the analysis
- The **analysis** object performs a piece of the analysis at every step of the traversal, accumulating the entire **analysis**



- Pick up a “clue” visiting each step of the **iterators** traversal, piecing together a map to the treasure, or final **analysis**



Spell Check Example



How do we do this?

- How does the analysis object distinguish different types of glyphs?
 - If/Else trees are ugly and nobody wants to deal with them
 - Switch statements aren't any better
 - Type checking would need to be updated every time a change was made to the glyph class hierarchy
 - Type checking is not effective in C++ and many other languages

How do we do this?

- Consider the following abstract operation added to the Glyph class:

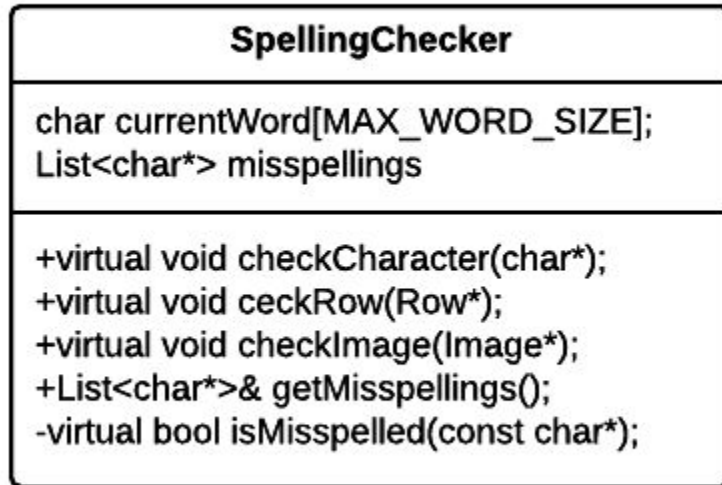
```
Void CheckMe (SpellingChecker&)
```

- Every Glyph subclass would define CheckMe

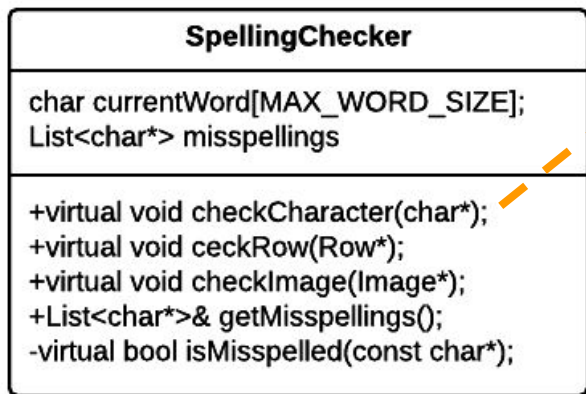
```
Void GlyphSubclass::CheckMe (SpellingChecker& checker) {  
    checker.CheckGlyphSubclass (this);  
}
```

- Here GlyphSubclass would be replaced by the name of each glyph subclass
- When CheckMe is called, the specific subclass is now known

How do we do this?



How do we do this?

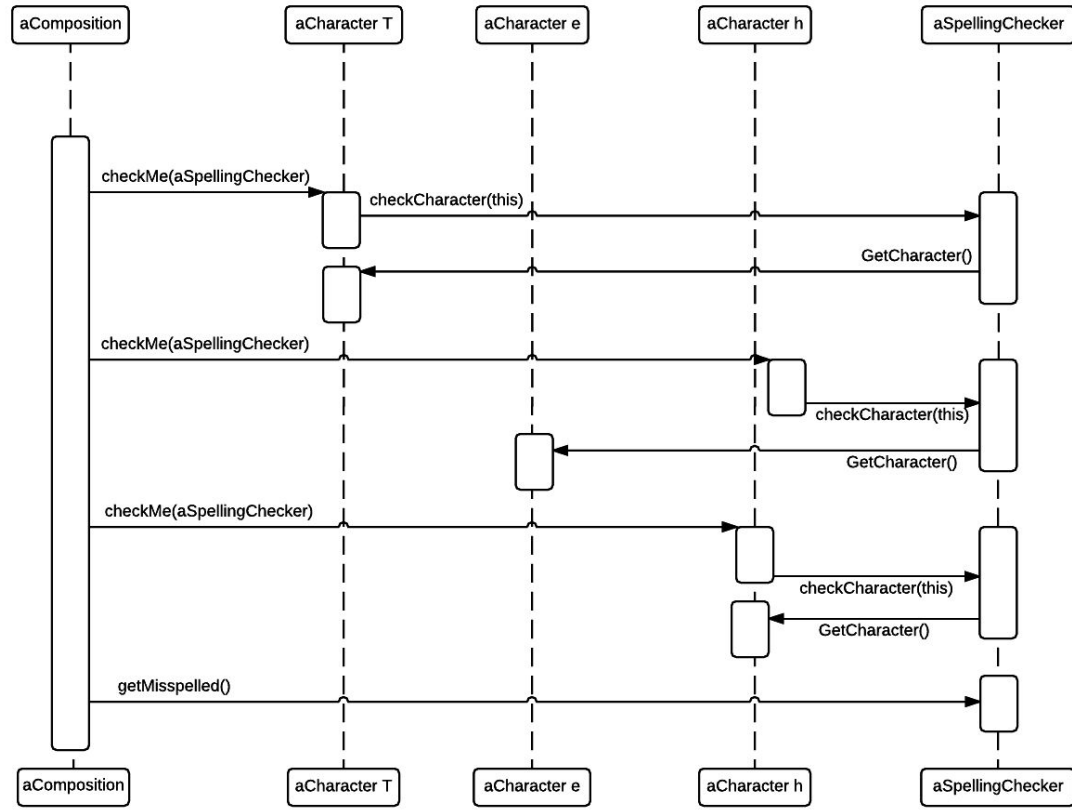


```
Void SpellingChecker::checkCharacter(char* c) {
    if(isalpha(c)) {
        //append to this->currentWord
    } else {
        // We hit a non-alphabetic character
        If (isMisspelled(this->currentWord) {
            // add to this->misspellings
        }
        // Reset the current word
        this->currentWord[0] = '\\0';
    }
}
```

How do we do this?

- Now we can traverse the glyph structure calling `checkMe` on each glyph with the `SpellingChecker` as an argument
- This identifies each glyph to the `SpellingChecker` and prompts the checker to do the next increment in the spell check

```
SpellingChecker* spellCheck;  
Competition* c;  
// ...  
PreorderIterator iter(c);  
For (iter.First(); !iter.is_done(); iter.Next()) {  
    Glyph* g = iter.Current();  
    g->checkMe(spellCheck);  
}
```

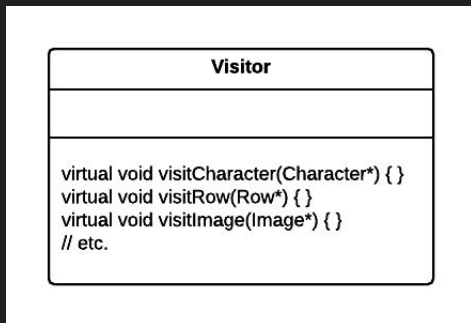


Supporting Multiple Analyses

- We can now find spelling errors using this approach
- However, this instance is specific to spelling error analysis (why?)
- By converting it to a more generic visitor pattern, we can extend the system to allow for different types of analysis easily
- The visitor interface `accept(Visitor*)` replaces the analysis-dependent operations like `checkMe(SpellingChecker&)`
- Specific visitor functions like `checkCharacter(char*)` are replaced by functions relating to a specific class `visitCharacter(Character*)`
- Different analyses can now inherit from this base class

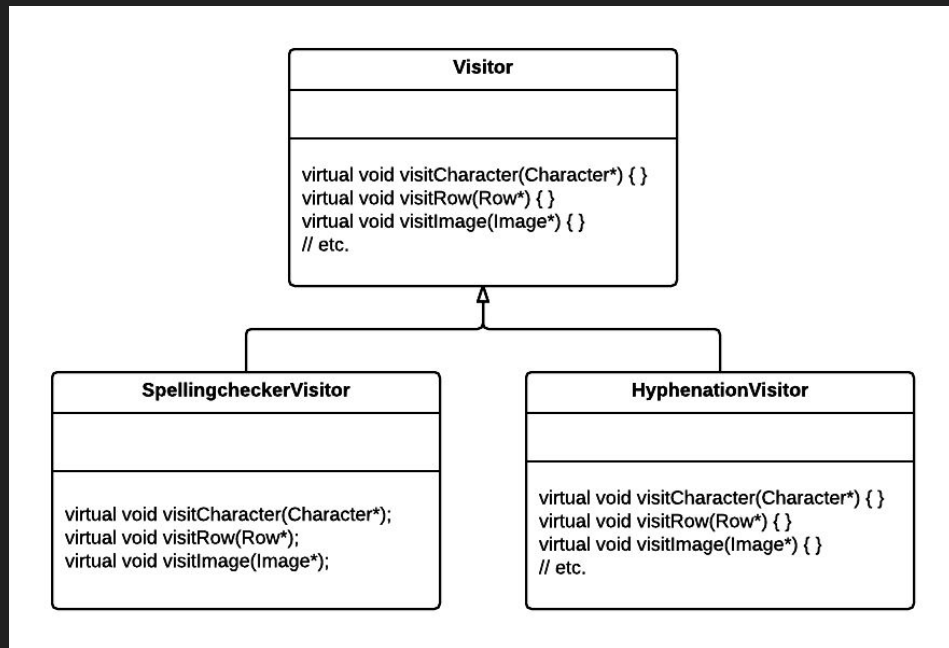
Visitor Class

```
class visitor {  
    public:  
        virtual void visitCharacter(Character*) { }  
        virtual void visitRow(Row*) { }  
        virtual void visitImage(Image*) { }  
        // Visit functions for all Glyph classes  
};
```

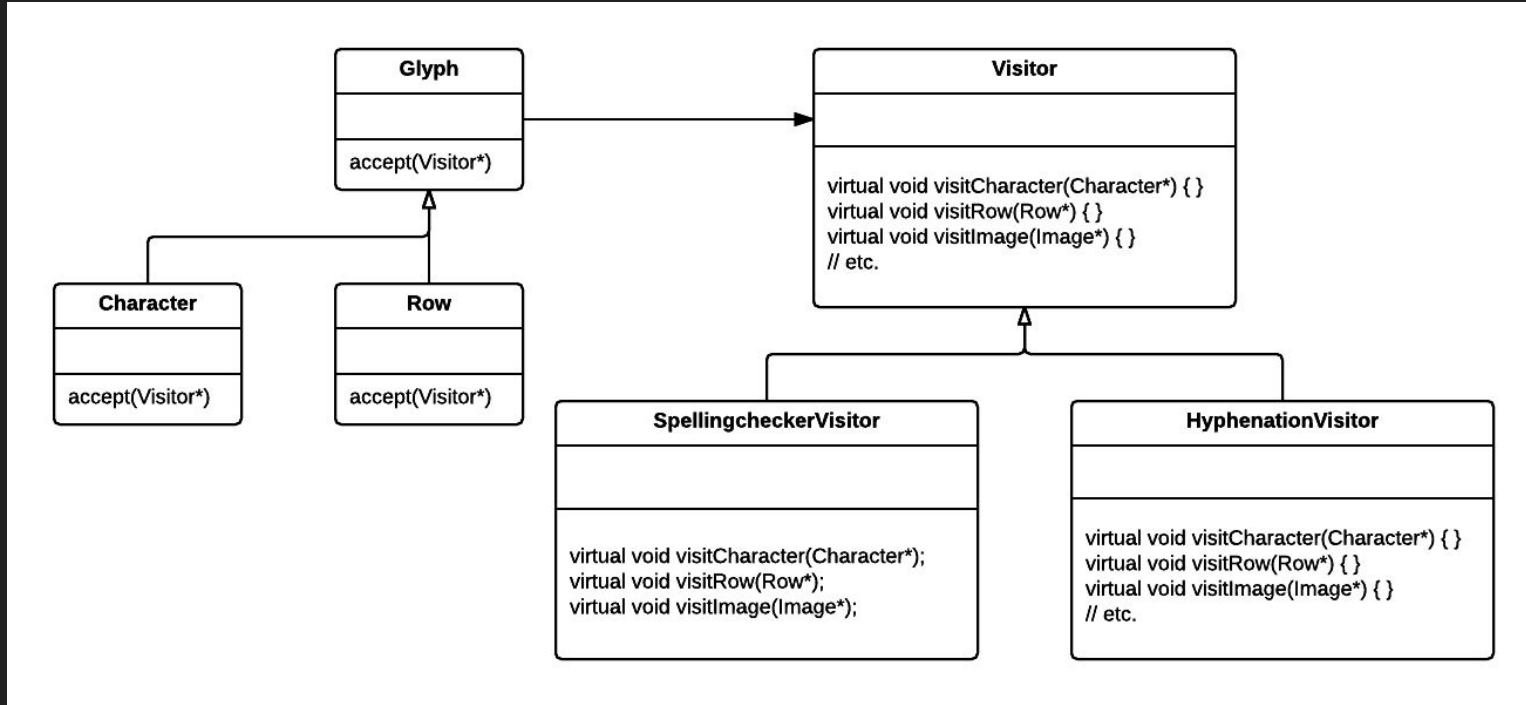


Visitor class and Concrete Subclasses

- `SpellingCheckerVisitor` is implemented the same way as before except `checkCharacter` is replaced with `visitCharacter`, etc.



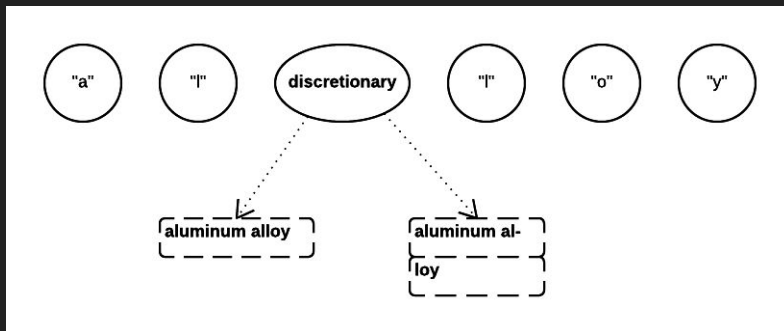
Visitor Class and Concrete Subclasses



How would we design a hyphenation
class as a visitor subclass?

HyphenationVisitor

- Traverse the text to accumulate characters into words
- Once an entire word has been assembled, apply an algorithm to determine potential hyphenation points
- At each hyphenation point, insert a **discretionary** glyph into the composition
- A discretionary glyph has two possible appearances, depending on whether or not it is the last character on a line, a hyphen or no appearance



Visitor Pattern Review

- The **Visitor** class and its subclasses allows an open-ended number of analyses on glyph structures without having to edit the glyph classes themselves
- Visitors can be applied to *any* object structure, not just composite structures
 - The object structure should have a traversal object (**iterator**)
- The visitor pattern is most suitable when you want to be able to do a variety of different things to objects with a stable class structure
- Adding a new kind of visitor requires no change to that class structure
- Adding a subclass to the structure requires an update to **all** visitor interfaces to include a `visit` operation for that subclass